

Twitter Sentiment Analysis

In this hands-on project, we will train a Naive Bayes classifier to predict sentiment from thousands of Twitter tweets. This project could be practically used by any company with social media presence to automatically predict customer's sentiment (i.e.: whether their customers are happy or not). The process could be done automatically without having humans manually review thousands of tweets and customer reviews.

1. Import libraries and datasets
2. Perform Exploratory Data Analysis
3. Perform data cleaning - removing punctuation
4. Perform data cleaning - remove stop words
5. Perform Count Vectorization (Tokenization) and TFIDF Vectorization
6. Build the suitable algorithms
7. Assess trained model performance